

Physically Based Rendering

Daichi Jameson

Sam Sartor

May 2017

1 Abstract

Physically based rendering (PBR) is a branch of computer graphics focused on the realistic and accurate description and simulation of materials' visual properties. Instead of describing the appearance of an object, an artist using PBR techniques endeavors to describe the material of the object itself. How metallic it is, how rough, how colored. In addition to an added degree of realism, PBR provides two additional benefits: a more intuitive method of manipulating materials and material consistency between different lighting environments [7]. While it is physically based, realism is not a requirement. PBR is also useful for cartoony or abstract visuals.

As a programmer, implementing physically based shading is challenging because of the complexity and obscurity of physically accurate lighting equations. Many are most commonly represented as massive integrals that are difficult to convert into usable code. Our PBR implementation uses a metalness workflow with the Cook-Torrance specular reflectance model, high dynamic range, normal mapping, shadow mapping, and gamma correction. The resulting program is able to create realistic and interesting images in real time.

2 Introduction

Developing intuitive and effective rendering techniques is a problem that has existed since the dawn of 3D graphics. Several different lighting models have come about in an at-

tempt to solve this problem, with varying degrees of success. One of the most famous lighting models is the Phong shading model, developed around 1975 by Bui Tuong Phong [4]. While the Phong shading model is able to create good looking scenery, it fails to account for certain real world laws, causing certain scenes to have unrealistic properties. More modern models account for additional real world properties, allowing for more realistic renders to be made.

Physically based rendering is a broad term describing a set of rendering rules that aims to emulate properties of the physical world in order to create realistic environments that are intuitive to manipulate. There is no limit on how pedantic PBR is allowed to be, but there are a few properties common to almost all PBR implementations: conservation of energy, the Fresnel effect, the microfacet surface model, and differentiating between metallic and dielectric materials [8]. In addition to the rules that are generally associated with PBR, several other rendering techniques are often used to maximize the realism of the final images.

3 Related Work

Physically based rendering aims to correct the consistency, realism, and controllability issues with older techniques like Phong. The first correction made by physically based rendering is conservation of energy. In theory, this is a simple physics principle: the amount of energy in a system must stay constant [1]. In the context of computer graphics, conser-




Figure 1: Reflecting microfacets




Figure 2: Reflection and scattering

vation of energy means that the light energy leaving a surface must be less than or equal to the light energy received by the surface (excluding emitted light) [7]. The next important concept to understand is the microfacet surface model (Figure 1). Most of the optical properties associated with real world surfaces originate with the tiny imperfections in the surface. A model of these imperfections, when used to develop improved shading equations, is called a microfacet model [2]. The Fresnel equation, nearly ubiquitous in modern computer graphics, is an example of a formula partially derived from microfacet models. The Fresnel equation is most commonly used to model the increased reflectivity of materials at shallow angles [7].

Perhaps the biggest correction made by physically based rendering is the distinction between metallic and non-metallic materials. To understand this distinction, it is important to first understand the distinction between specular (reflected) and diffuse (scattered) light. When photons hit a surface,

some reflect immediately off the surface and back into free space. This is the specular component. The “hardness” of this reflection is due only to the roughness of the tiny imperfections on the surface. Hence the usefulness of microfacet models. The remainder of the photons can travel past the first layer of atoms and into the volume of the material itself (all materials are somewhat transparent). Here they are scattered and reflected many times, moving in an irregular path through the object (Figure 2). Many of these photons will return to the surface and travel off in a random direction. This “scattering” creates the diffuse shading component. Keep in mind that in both these cases, some of the photons will be absorbed as heat before scattering or reflecting.

Non-metallic (dielectric) materials happen to reflect all wavelengths evenly and so have a white specular component. Color is only added during scattering. The ratio between reflected and scattered light is described primarily by the Fresnel equation. On the other

hand, any light that passes the surface of a metallic material will never return. Thus, metals only have a specular component. Diffuse looking metal objects are simply rough, not truly diffuse. Unlike dielectric objects, the ratio of reflected light is wavelength dependent. This causes the specular component of metals to be colored [2].

There are two primary workflows used to describe the properties of a material: the specular workflow and metalness workflow [3]. The metalness workflow aims to provide an intuitive method of describing materials. The inputs to this workflow are albedo, metalness, and roughness [3]. Albedo describes the color of the material. This is the diffuse color of a dielectric and the specular color of a metal [3]. The metalness map then describes which of those cases applies [3]. The roughness map describes the roughness of the material, which affects the sharpness of specular reflections[3]. Conversely, the specular workflow aims to provide inputs similar to Phong, where materials are described using their explicit color values [3]. Specular workflows use albedo, specular, and glossiness maps [3]. An important difference between Phong and the specular workflow is that the specular workflow keeps all of the color information in the albedo map, and uses the specular map to only relay information on how specular a surface should be [3]. Regardless of the workflow or implementation, these inputs are almost always stored in a texture to allow for variation across an object.

In addition to the basic principles of physically based rendering, there are a few other rendering techniques that allow for more re-




Figure 3: Inputs to metalness workflow

alistic looking scenery. The use of high dynamic range (HDR) and gamma correction is particularly important to achieving realism and/or consistency. A common issue in normal renders occurs when the final light value of a pixel exceeds 1.0. Lighting information about the scene is then lost, and the final render becomes blown out. High dynamic range addresses this issue by first rendering the scene on an unbounded scale, and afterwards transforming the lighting values into the monitor’s display range through a process called tone mapping[6]. Tone mapping allows for the lighting to be scaled down into a displayable range while still having control of the final image. Gamma correction is another extremely important rendering technique that converts lighting values between a

linear color space and the color space monitors use to render images [5]. Using gamma correction ensures that the values calculated in shaders are physically correct, something that is critical for physically based rendering to be done correctly.

4 Problem Statement

Creating realistic renders in real time is a challenging task. While it is currently (and may always be) impossible to simulate the world perfectly, there are several techniques that attempt to emulate real world properties in a computationally efficient manner whilst still preserving render quality. In addition, it is important for materials to be easily adjustable, and for materials to stay consistent between different lighting environments.

The Phong model is one of the most popular solutions for real time rendering. Though it is a fairly effective solution, it has many shortcomings and does not account for a few key physical properties, which limits the degree of realism it can achieve. Another major disadvantage of the Phong lighting model is the difficulty of working with material colors. The Phong lighting model requires the user to directly manipulate the specular, diffuse, and ambient values, which can be quite difficult to adjust in order to create realistic looking materials. Additionally, the Phong lighting model suffers from consistency issues on materials between different lighting environments [7]. In order to address this issue, an artist would have to go and tweak the colors of a model to make sure that its look re-

mains consistent [7]. Physically based rendering aims to solve these issues.

5 Problem Solution

We implemented this project using a language called Rust. Rust is more robust and higher level than C++ while preserving on-the-metal performance and access. We hoped that using Rust would make our implementation easier to code and more robust. High-level, reliable, and multi-threadable access to OpenGL is provided through the gfx library. In addition to OpenGL, gfx can use DirectX 11, DirectX 12, Metal, and Vulkan as backends; however, this project exclusively uses the OpenGL backend.

We chose to use the metalness workflow for several reasons. It seems to be better documented, we were able to find more test models using it, and it made more intuitive sense to us.

In addition to a basic physically based shading model, our program improves realism using HDR rendering, gamma correction, normal mapping, and shadow mapping.

On startup, the program loads a set of objects. Each one comprises a Wavefront OBJ model and 4 PNG images (normal, albedo, metalness, and roughness). Position, normal, and texture coordinates are loaded from the OBJ file and put in vertex buffer along with calculated tangent and bitangent vectors. Triangle indices are also loaded from the file and placed into an index buffer. Similarly, the 4 images are loaded into opengl textures. Only the albedo texture is in sRGB

space and needs to be corrected, the others are linear (note that this is a common convention but not universal).

The rendering step uses multi-pass techniques extensively. The first pass renders the scene to a gbuffer, which stores position, UV, and normal information. This step is also responsible for applying the object's normal map. Applying the normal map early prevents us from having to store the tangent and bitangent in the gbuffer, reduces the number of textures being bound by the lighting pass, and reduces the amount of work being repeated by the lighting pass.

The next two passes are repeated per-light. First, the scene is rendered from the light's perspective in order to populate a shadow map. Shadow acne is avoided by clipping front faces and only storing the depth of back faces. This does not create any additional issues because all of our example objects are solid.

Once the shadow pass is complete for the given light, a fullscreen quad is rendered with the PBR fragment shader. This shader uses the Cook-Torrance model which is energy conserving, microfacet based, and includes the Fresnel effect [2]. This shader is also responsible for the shadow depth-check. A total of 6 textures are loaded: gbuffer layer a, gbuffer layer b, the albedo map, the metalness map, the roughness map, and the shadow map. This shader outputs a floating-point luminance value which is added to the luminance buffer (additive blending mode). The luminance buffer is reset every frame, and then the contribution from each light is added to the buffer in sequence. This al-

lows us to include multiple shadow-casting lights into our scene without loops or arrays in the fragment shader. By alternating between shadow pass and lighting pass, a single shadow map is required. It can be written over by every light because the previous light's contribution has already be added to the total luminance.

A final fullscreen quad is rendered to convert the 32-bit HDR luminance buffer into an 8-bit LDR screen buffer. This final pass is responsible for gamma correction and tone mapping.

6 Results

Our results are very encouraging. Using textures provided us a level of detail and storytelling ability that would be impossible to obtain otherwise. Rust and gfx also proved to be effective: adding new passes, buffers, shaders, textures, etc. was clean and straightforward. The program also performed well, handling 10 or more lights (each with shadow passes) without a noticeable frame-rate drop. Most importantly, our objects appear consistent under different lighting conditions. Wood looks like wood, copper looks like copper, and steel looks like steel, regardless of what color lights we choose. If we had the time, additional improvements could be made by adding ambient occlusion, antialiasing, and/or image-based lighting.



Figure 4: Car rendered in our program

7 Conclusion

Throughout the history of computer graphics there have been many attempts to model the lighting and shading of scenes. The Phong model has become one of the most prominent and visually appealing shading model. Despite the Phong model’s popularity, it fails to account for several properties of the real world, such as the difference between metallic and dielectric materials, the increase in specularities at steep angles, and the conservation of energy. Physically based rendering is a new paradigm that appeared and declared that the end result is not the only thing that matters, but rather that properly modeling real world phenomena will lead to better looking, more physically accurate results. While modeling the world as accurately as possible leads to more complex math and more computationally intensive calculations, new hardware has given us the ability to do these calculation in real-time.

Our implementation, while not trivial to create, does a good job demonstrating these advantages. Through a collection of multipass techniques, we can render a variety of complex and visually appealing objects. These advantages are perhaps why almost all modern game engines and visual effects use physically based rendering.

References

- [1] *Conservation of Energy*. URL: <https://www.grc.nasa.gov/www/k-12/airplane/therm01f.html>.
- [2] Robert L. Cook and Kenneth E. Torrance. “A reflectance model for computer graphics”. In: *ACM SIGGRAPH Computer Graphics* 15.3 (Jan. 1981), pp. 307–316. DOI: 10.1145/965161.806819.
- [3] *PBR Texture Conversion – Marmoset*. URL: <https://www.marmoset.co/posts/pbr-texture-conversion/>.
- [4] Bui Tuong Phong. “Illumination for computer generated pictures”. In: *Communications of the ACM* 18.6 (Jan. 1975), pp. 311–317. DOI: 10.1145/360825.360839.
- [5] Joey de Vries. *Gamma Correction*. URL: <https://learnopengl.com/#!Advanced-Lighting/Gamma-Correction>.
- [6] Joey de Vries. *HDR*. URL: <https://learnopengl.com/#!Advanced-Lighting/HDR>.

- [7] Joey de Vries. *PBR Theory*. URL: <https://learnopengl.com/#!PBR/Theory>.
- [8] Brian Yu. *Physically Based Rendering Encyclopedia*. URL: https://docs.google.com/document/d/1Fb9_KgCo0noxROKN4iT8ntTbx913e-t4Wc2nMRWPzNk/edit#heading=h.y94w1qp4r4mt.